

1 Intro to trees

1.1 ML Pipeline

- 1.1.1 Definition of Models
 - Model can refer to equations (linear/non-linear), code/rules, Neural networks, decision trees, bayesian graph, Deep NN
 - In simple terms, model defines a function that maps inputs to outputs and is mostly mathematical equations

1.1.2 Machine Learning Pipeline

Machine learning follows an iterative pipeline in the following order:

Data collection → Data Extraction → Data pre-processing → Model choice/design → Model training → Model validation (evaluation) → Model understanding (explainability) → Model deployment

2 Paradigms of ML

2.1 Classes of ML techniques

- 1. Supervised learning
 - Given a set of data features and labels, design an algorithm that predicts label for new data (never observed before).
 - Regression task: e.g. predict price (continuous value) of houses given existing data label (house size)
 - Classification task (binary): e.g. predict (discrete value) whether benign or malignant tumor based on tumor size
 - Classification task (multiple classes): e.g. predict class of disease given size of tumor
 - Classification task (multi-feature)
- 2. Unsupervised learning
 - In supervised learning, data and label always comes in a pair but in unsupervised learning, it is likely that we do not have label or any additional information
 - The goal is to find structure in data that solves tasks like clustering, classification, compression and generative model
- 3. Reinforcement learning (not tested)
 - Learns a sequence of actions that maximizes a reward
 - RL agents learns to plan the future to win
 - Main issue is that lots of data is required because reward is sparse (e.g. whole chess game must be played to get a score of +1, -1 or 0)
- 4. Self-supervised learning (not tested)

2.2 Supervised Learning with kNN

- Goal is to estimate a function f such that $y = f_{\theta}(x)$
 - x is a raw data point
 - y is the label (either a real number for regression task or categorical variable for classification task)
 - θ are the parameters of the predictive function
 - y is annotated by humans which is extremely time consuming, expensive, biased and is limited to human knowledge

2.2.1 Difference between kNN, Decision tree, SVM

- What differentiates these three models?
 - Geometry of decision function/boundary
 - Learnable vs non-learnable models
 - Complexity (speed and memory) of training and inference

2.2.2 k-Nearest-Neighbor

- Assumes that close data points have similar labels
- Basic Algorithm to predict the label of a new data point x^* :
 - Find the closest k data points in the training set, $S = \{(x_i, y_i)_{i=1:N}\}$
 - Assign the label of the **mean/mode** to the new data point
- Distance metric:
 - L2/Euclidean distance
 - L1/Manhattan distance
 - Cosine distance (angle)
 - Jaccard distance (sets)
- Larger the value of k , the more smooth the decision boundary, less overfitting
- Hyper-parameter k acts as regularizer which increases robustness (new & similar input used, output is same) of predictor
- Terrible at high-dimensionality since data points sampled from random distribution will have about the same distance from each other (curse of dimensionality)
- Performs well for few number of meaningful features
- Can be used for both regression and classification tasks

2.2.3 Time Complexity of kNN

- No training needed
- $O(n \cdot d)$ for nearest neighbor and $O(n \cdot d \cdot k)$ for kNN
 - n = number of training data points
 - d = number of data features
 - k = number of nearest neighbor
- High memory usage as all training data is loaded into memory
- k-d tree and hashing techniques to speed up

2.2.4 Curse of Dimensionality vs Blessing of Structure

- Curse of dimensionality states that if points were chosen randomly, distance between them are about equal → kNN will not work well
- Blessing of structure instead says that **real world data is unlikely to be randomly distributed** and have structures like edges or textures (for images)
 - This brings down dimensionality of data to be $< R^d$

2.2.5 Pros and Cons of kNN

- Extremely simple and expressive (can produce non-linear boundary decision)
- As $n \rightarrow \infty$, kNN is provably very accurate but requires huge amount of space
- As $d \rightarrow \infty$, kNN fails due to curse of dimensionality

3 Decision Trees

3.1 Motivation of kd-tree

- kNN is extremely slow and time consuming when we use a large n since we need to compute distance to all of the points each time we make a query
- Want a way to **prune away points** that we know are extremely far away to cut computation time

How kd-tree works

3.2.1 Tree Construction

- Split recursively in half along each feature dimension
 - We want to have about the same data points at each side of the split for max efficiency
 - Which direction we choose to split first is chosen greedily
- Iterate over all feature dimensions
- End up with a depth of $O(\log_2 n)$

* In general, number of partition is about 2^d to capture largest variation of data

3.3 Searching k-d tree

- Given a point x , identify which set X lies in and find the closest neighbor in that set, x_{NN}
- Compute distance $d(x, C)$, where C is the cut
- If $d(x, C) > d(x, x_{NN})$, discard everything that is in the set of the cut, else find nearest neighbor in the other set of the cut
- Repeat for all cuts

3.4 Complexity of k-d tree

- Space complexity of k-d tree: $O(2^p) \rightarrow O(n)$, where $p = O(\log_2 n)$
- Time complexity to build k-d tree: $O(n \log_2 n)$
- Inference speed:
 - Best case: $O(\log_2 n + d)$
 - Worst case: $O(dn)$, basically nearest neighbor
 - Average case: $O(d \log n)$

3.5 Pros and cons of k-d tree

- Pros:
 - Exact kNN but no need to backtrack in parent nodes
 - Easy to implement
 - Average inference much faster than kNN, $O(d \log_2 n)$ vs $O(dn)$
- Cons:
 - Cuts are axis aligned → not good in high dimensionality

3.6 Decision trees

- kNN requires full training data to make prediction
- k-d tree uses the fact that data concentrate in regions to speed things up
- Ultimately, goal is not to find closest data points but to get a classification or regression value
 - e.g. If new data point falls in cluster of 1000 +ve points, can just classify as +ve without calculating distance
 - Want to leverage the idea of clustering so that no need to load full training set

3.6.1 Inference using Decision Tree

- Once tree constructed, don't need to keep training set in memory
- Only need to store **tree structure** ($O(\log_2 n)$ depth) as well as **class probability/regression value/class label** in leaf nodes
- Inference is extremely fast at $O(\log_2 n)$, independent of d

3.6.2 Optimal decision tree


- In theory if no 2 data points have same features but different labels, we can ensure min depth
- In practice, finding minimum size tree is NP-hard
- Compromise is to **minimize a function that measures label purity**

3.6.3 Purity

- Purity is defined as the fraction of data with label k in a set, $p_k = \frac{|S_k|}{|S|}$
- Worst case is when all leaves are random prediction, i.e. $p_k = \frac{1}{c} \cdot \forall k$, where c is the number of classes
- Want to maximize Kullback-Leiber distance between random prediction and best candidate p

3.6.4 Entropy

- Maximizing KL distance is equivalent to minimizing entropy
- If probability of each class is equal, then entropy is at its maximum



$$H(L, R) = p(L)H(L) + p(R)H(R) \\ = \frac{|L|}{|S|}H(L) + \frac{|R|}{|S|}H(R)$$

- Entropy $H = -\sum p_k \log p_k$
- Ideally, entropy of L and R should be less than S, i.e. $H(L), H(R) < H(S)$

3.6.5 Information Gain

- Information gain (IG) is the difference between the entropy of the original set S and the weighted sum of the entropy of the subset S_k


$$IG(S, S_1, \dots, S_c) = H(S) - H(S_1, \dots, S_c) = H(S) - \sum_{k=1}^c p(S_k)H(S_k)$$

3.6.6 Construction of Decision Tree

- Goal is to find subsets that maximizes the information gain, achieving the purest possible subsets.
 - Identifying the purest subsets is to find a feature x_i and a threshold value t_i .
- Decision tree construction (pseudo-code)
 - While leaf nodes are not pure (or \geq threshold)
 - Loop over (remaining) feature dimensions, i.e. x_1, x_2, \dots, x_d
 - Loop over n thresholds (e.g. middle points between two consecutive points, such as $t_i = (x_{i-1} + x_i)/2$)
 - Compute information gain for R and L.
 - Save (dimension, threshold value) with maximum information gain.
 - Split space with best (dimension, threshold value) and remove dimension x_i from loop.
- Exact complexity is $O(n \cdot d)$. Approximations are used in practice for speed-up.

3.7 Regression trees

- It is easy to extend decision trees to regression trees as long as a purity function can be defined for the new task



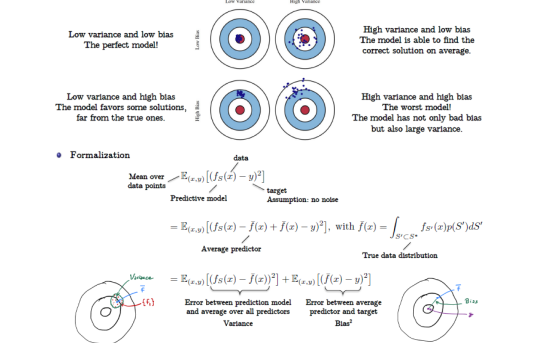
3.7.1 Complexity of Regression Tree

- Regression tree space, time and inference complexity all follows that of k-d tree

3.8 Bagging

- Decision trees are extremely fast but have **high variance**, i.e. weak learners
- High variance**: quality of the classification/regression solutions vary significantly

3.8.1 Bias and Variance



3.8.2 Reducing Variance

- Most common way to reduce variance is to take average of several solution, i.e. **ensembling**
- Relies on the **law of large numbers** where the average of many small predictors will tend to the average of the true predictor
- However this requires us to have access to more training sets which in most cases we do not have

3.8.3 Bagging algorithm

- Sample m datasets S_1, \dots, S_m from original S with replacement
- For each training set S_j , train a classifier f_{S_j}
- Final ensemble classifier is $\hat{f}(x) = \frac{1}{m} \sum_{j=1}^m f_{S_j}(x)$

* Note that **sampling with replacement breaks the assumption of independent and identically distributed** (i.i.d) data and therefore does not have **theoretical guarantee** that the ensemble will give a good estimation of the true mean actual predictor

3.8.4 Advantages of bagging

- In practice, bagging reduces variance quite effectively, but after a large number m , will have diminishing returns
- Bagging can **reduce the variance without increasing the error** of an unbiased classifier

3.8.5 Random forests

- One of the most popular and useful bagging algorithms
- Random forest is an ensemble of decision trees
- Hyper-parameters: $k = \sqrt{d}$

3.8.6 Boosting

- Boosting is used to help **reduce bias** on weak learners such as decision trees with limited depth
- Uses the idea of ensembling which combines large number of weak learners to generate a strong learner with low bias
- Note that both **boosting and bagging are both agnostic to algorithm used**

4 Linear Models

4.1 Applications of Linear Models

- Classification: $\text{sign}(\theta^T x)$, $y = \pm 1$, e.g. Approve or reject
- Regression: $\theta^T x$, $y \in \mathbb{R}$, e.g. Amount of credit
- Logistic regression: sigmoid($\theta^T x$), $y \in [0, 1]$, e.g. probability of defaulting

4.2 Equation of Linear Model

- $f_{\theta}(x) = \theta^T x = \sum_{i=1}^d \theta_i x_i$, where θ_i is the weight/importance of feature x_i

- Linear classification: $\begin{cases} \theta^T x > 0 \Rightarrow +1 \\ \theta^T x < 0 \Rightarrow -1 \end{cases}$

- Linear regression: $\theta^T x \Rightarrow$ Amount (scalar)

4.3 Loss function

$$L(\theta) = \frac{1}{n} \sum_{j=1}^n (f_{\theta}(x^j) - y^j)^2 - \text{Mean squared error}$$

where $f_{\theta}(x^j)$ is the predicted value from the model and y^j is the actual labeled value

- Goal of models is to choose a θ that **minimizes** the mean squared error


4.4 Finding best θ

- Mainly uses 2 approaches: **Normal equations** and **Gradient descent**
- Normal equation solution (**only for linear regression**):
$$\theta = ((X^T X)^{-1} X^T) y$$

where $(X^T X)^{-1} X^T$ is known as the pseudo-inverse of X and is **computationally expensive** to compute if n or m is large

4.4.1 Gradient Descent

- Works in high-dimensional spaces, convergence **independent of data dimensionality d**
- Steps to carry out gradient descent
 - Start at a random $\theta(t)$
 - At each step, update the value of parameter: $\theta(t+1) = \theta(t) + v$, where v is the opposite direction of the greatest +ve gradient
 - Stop when gradient $<$ arbitrary threshold
- $v = -\eta \frac{\partial L}{\partial \theta}(\theta(t))$, where η is the step size
 - No good way to find the right step size, just have to try a few values and pick best on



4.4.2 Stochastic Gradient Descent

- Main problem with gradient descent is that there are **lots of saddle points / flat regions** in **high dimensional spaces**

Stochastic gradient descent

- A fast variation of gradient descent that considers only a small set of data points to update the value of the parameters is stochastic gradient descent (SGD):
 - Pick a (random) small subset of m training data (e.g. a single/512 data points), i.e. $(x^{(k)}, y^{(k)})$
 - Compute the loss value for this subset, i.e. $L = \frac{1}{m} \sum_{i=1}^m L(f_{\theta}(x^{(k)}), y^{(k)})$
 - Compute the gradient for this subset, i.e. $v = -\frac{1}{m} \sum_{i=1}^m \nabla L(f_{\theta}(x^{(k)}), y^{(k)})$
 - Update parameters, i.e. $\theta(t+1) = \theta(t) - \eta v$
- Advantages
 - Faster update of parameters : Gradient computed with $m=512$ rather than all data points.
 - Stochastic optimization : Helps escape saddle points in high-dimensional space
 - Simple to implement

4.4.3 Gradient Descent vs Normal Equation

- Two methods for solving the linear regression task :
 - Gradient Descent
 - Works well for very large n , training data, with SGD (300B tokens with GPT3)
 - Works well for very large d , #data features ($d=1M$ with 1,000-1,000 images)
 - Works well for very large $|\theta|$, #parameters features (175B with GPT3)
 - Very slow and requires to select time step η
 - Normal Equations
 - Very fast for $n=O(10^6)$ data points, do not require to choose η
 - Need to compute $(X^T X)^{-1}$, $O(n^3)$ operation but faster approximations exist
- Gradient Descent is a universal optimization technique as long as the considered loss is continuous and differentiable (as gradient is required).
 - GD does not work for discrete losses like win/lose at the game of Go.

4.5 Logistic Regression


- Used to **predict probability** e.g. probability of heart attack, given training data of person's health information and whether they have a heart attack
- $P_{\theta}(y|x) = \begin{cases} f_{\theta}(x) & \text{for } y = +1 \\ 1 - f_{\theta}(x) & \text{for } y = -1 \end{cases}$
- Goal is to minimize cross-entropy loss (a.k.a. log loss):

$$\min_{\theta} L(\theta) = \frac{1}{n} \sum_{j=1}^n \log(1 + \exp(-y^j \theta^T x^j))$$

where n is the number of training data

4.6 Support Vector Machine (SVM)

- Goal of SVM is to find a linear separator that partitions the feature space into 2 regions

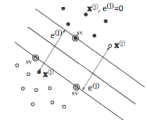


- We want margin d to be as large as possible to generalize to more data points
- $d = \frac{2}{\|\theta\|}$

- 4.6.1 **Soft-margin SVM**
- Aims to solve the problem that real-world data is typically **non-linearly separable** due to outliers \rightarrow no mathematical solution to SVM
 - Idea is to **introduce a slack variable $e(f)$** for each data point that represents the prediction error

$$\min_{\theta} \|\theta\|^2 \text{ such that } \begin{cases} f_{\theta}(x) \geq 1 \text{ for } x \in R_+ \\ f_{\theta}(x) \leq -1 \text{ for } x \in R_- \end{cases} \quad \text{Standard SVM}$$

$$\Downarrow$$

$$\min_{\theta, e} \|\theta\|^2 + C \cdot \sum_{i=1}^n e^{(i)} \text{ such that } \begin{cases} f_{\theta}(x^{(i)}) \geq 1 - e^{(i)} \text{ for } x \in R_+ \\ f_{\theta}(x^{(i)}) \leq -1 + e^{(i)} \text{ for } x \in R_- \\ e^{(i)} \geq 0, C \geq 0 \end{cases} \quad \text{Soft-margin SVM}$$


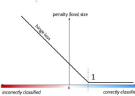
- Soft-margin **penalizes misclassifications and correct classifications that fall within margin**
- Typically would have a **much smaller margin** than standard SVM
- Uses hinge loss function:

$$\min_{\theta, e} \|\theta\|^2 + C \cdot \sum_{i=1}^n e^{(i)} \text{ such that } \begin{cases} f_{\theta}(x^{(i)}) \geq 1 - e^{(i)} \text{ for } x \in R_+ \\ f_{\theta}(x^{(i)}) \leq -1 + e^{(i)} \text{ for } x \in R_- \\ e^{(i)} \geq 0, C \geq 0 \end{cases}$$

$$\Updownarrow$$


$$\min_{\theta} \|\theta\|^2 + C \cdot \sum_{i=1}^n \max(0, 1 - y^{(i)} f_{\theta}(x^{(i)}))$$

Hinge loss



4.7 Kernel SVM

- Linear models are limited to **linearly separable data points** \rightarrow cannot handle complex/non-linear data sets
- To use a linear kernel, one must first map data from original space R^d to a higher dimensional space $R^b, b \gg d$ so that the classes can be distinguished using linear functions



5 Bias & Variance

5.1 Definition of Bias, Variance and Noise

- Bias: **intrinsic error/difference** between the average prediction model and the true regression value. **Inherent to the model and independent of data**. Also known as **model error**

- Variance: captures how much the learner changes when it is computed on different training sets. Indicates **expressivity** \rightarrow complex algo have high expressivity

- Noise: ambiguity inherent to the data distribution and feature representation. **Impossible to get rid of** and is often modeled as a stochastic process that is added to "clean" data
- Good hypothesis has low bias and variance
 - Compare between 2 models, just add their bias and variance and compare the sum, lower the better

5.2 Bias-Variance-Error Decomposition


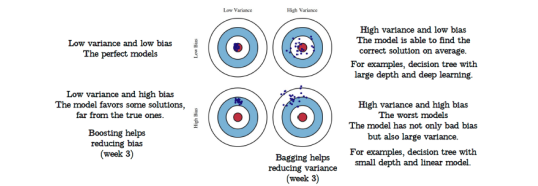
$$\underbrace{\mathbb{E}_{x,D}[f_D(x) - y]^2}_{\text{Expected test error}} = \underbrace{\mathbb{E}_{x,D}[f_D(x) - f(x)]^2}_{\text{Variance}} + \underbrace{\mathbb{E}_x[(f(x) - y(x))^2]}_{\text{Bias?}} + \underbrace{\mathbb{E}_{x,D}[(y(x) - y)^2]}_{\text{Data noise}}$$

- Most fundamental equation of supervised learning known as the **bias-variance-error trade off**

5.3 Over/Under Fitting

- Over-fitting: Occurs when learner $f_D(x)$ has zero prediction error on a training set D , but have large test error
 - Learner is too expressive and will become over-specialized of the training data, unable to extrapolate to unseen data.
 - Typically have high variance to capture complexity of whole dataset
- Under-fitting: learner $f_D(x)$ which is not able to predict correctly a training set D (very unlikely to happen), have both **high training and test error**
 - Learner is not expressive enough. It will make error on the provided training set, i.e. unable to benefit from all information present in the training data
 - Typically have no enough variance/complexity \rightarrow just using a simple algorithm

5.4 Bias-Variance Trade-off



5.5 How to reduce variance and bias

- How to reduce high variance?
 - Reduce model complexity : **Lower model expressivity** or use regularizers (week 6)
 - Remove non-informative/bad data features : E.g. house features such house color, back door, etc but challenging to decide bad features (not done in practice).
 - Bagging (week 3) : Averaging weak high-variance learners produces strong learner with low variance (requires fast computation of weak learners in practice).
- How to reduce high bias?
 - Increase model complexity : **More expressive models** (deep learning)
 - Add more informative data features : E.g. house features such as storage space, garden, security system, etc (effective but time and money consuming).
 - Boosting (week 3) : Adding weak high-bias learners produces strong learner with low bias (requires fast computation of weak learners in practice s.a. small-depth decision trees).
- Add more training data or use data augmentation to reduce both bias and variance.

6 Overfitting Regularization

6.1 Loss Regularization

- Method to **reduce overfitting** by **decreasing expressivity** of model (e.g. from 10th order to 2nd order function)
- Typically use Regularization with stochastic gradient descent \rightarrow SGD speeds up gradient descent and also regularizes predictive function (w.r.t. θ), allowing for better generalization
- Uses the concept of **constraint optimization with soft constraints**, basically saying that we are setting the $\theta_3, \theta_4, \dots, \theta_n$ of a hypothesis space $\mathcal{H} = \theta_1 + \theta_2 x + \theta_3 x^2 + \dots + \theta_n x^{n-1}$ to $\leq C$

$$\min_{\theta} L_{\theta} \text{ such that } \sum_{j=0}^d \theta_j^2 = \theta^T \theta \leq C, C > 0 \quad (1)$$

- If C is small $\rightarrow \theta_j \geq 3 \approx 0$, i.e. $\mathcal{H}_{10} = \mathcal{H}_2$
- If C is large, then basically doing unconstrained optimization (MSE)

6.2 Relationship between constraint regularized problem and unconstrained optimization problem


- For every general constraint regularized problem like in (1), there exist and equivalent unconstrained optimization problem (**easier to solve**)

$$\min_{\theta} L(\theta) + \lambda \theta^T \theta, \lambda > 0 \quad (2)$$

- For each value C , there exists a value λ such that $(1) \equiv (2)$ (Lagrange multiplier), $C \propto \frac{1}{\lambda}$

6.3 Original MSE loss vs Regularized MSE

<p>Original MSE loss :</p> $\min_{\theta} L(\theta) = \frac{1}{n} \sum_{i=1}^n (\theta^T x^{(i)} - y^{(i)})^2$ <p>Set gradient of loss to zero :</p> $\nabla L = \frac{\partial L}{\partial \theta} = 0 = \frac{1}{n} X^T (X\theta - y) = 0$ $\Rightarrow \theta = (X^T X)^{-1} X^T y$	<p>Regularized MSE loss :</p> $\min_{\theta} L(\theta) = \frac{1}{n} \sum_{i=1}^n (\theta^T x^{(i)} - y^{(i)})^2 + \lambda \theta^T \theta \leq C$ $\Leftrightarrow \min_{\theta} L(\theta) = \frac{1}{n} \sum_{i=1}^n (\theta^T x^{(i)} - y^{(i)})^2 + \lambda \theta^T \theta$ <p>Set gradient of loss to zero :</p> $\nabla L = \frac{\partial L}{\partial \theta} = 0 = \frac{1}{n} X^T (X\theta - y) + \lambda \theta = 0$ $\Rightarrow \theta = (X^T X + \lambda I)^{-1} X^T y$
<p>Minimizer of the total loss :</p> $\theta^* = \arg \min_{\theta} L_{\text{total}}(\theta) \Leftrightarrow L_{\text{total}}(\theta) \leq C$ $\Leftrightarrow \theta^* = \arg \min_{\theta} L_{\text{total}}(\theta) + \lambda L_{\text{total}}(\theta)$	



6.4 L2 vs L1 regularization


- The L_2 -ball regularization can be generalized to L_p -ball, $p \in [0, +\infty]$.

$$\min_{\theta} L_{\text{total}}(\theta) \text{ s.t. } \|\theta\|_p \leq C \Leftrightarrow \min_{\theta} L_{\text{total}}(\theta) + \lambda \|\theta\|_p, \text{ where } \|\theta\|_p = \left(\sum_{i=1}^n |\theta_i|^p \right)^{\frac{1}{p}}$$
- L_p -ball/ L_p regularization, a.k.a. weight decay
 - Advantages : Strictly convex, differentiable, fast optimization, robust w.r.t. **perturbation**. This means no feature selection in e.g. L_0 ($L_0 = \theta_0 + \theta_1 x + \theta_2 x^2 + \theta_3 x^3 + \dots + \theta_n x^{n-1}$, as all data features are used for prediction).
 - Limitations : Although θ_i values are minimized, solutions are dense, i.e. $\theta_i > 0$. This means no feature selection in e.g. L_0 ($L_0 = \theta_0 + \theta_1 x + \theta_2 x^2 + \theta_3 x^3 + \dots + \theta_n x^{n-1}$, as all data features are used for prediction).
- L_1 regularization
 - Advantages : Convex (but not strictly), fast optimization algorithms exist, robust w.r.t. perturbation, solutions are **guaranteed to be sparse meaning feature selection**, as only a few data features are used for prediction.
 - Limitations : Not differentiable at the origin.

6.5 Cross Validation

- We need a surrogate of the test set to estimate the regularization parameter λ which identifies the right model complexity that minimizes the test error
- Since we cannot touch data from the test set, we **split training set into training and validation set**

- Use p hypothesis/values to estimate λ .



- Estimation of error of validation set $L_{\text{val}} = L_{\text{test}} \pm O\left(\frac{1}{\sqrt{m}}\right)$, where m is the number of data points in validation set
 - To have L_{val} be as close to L_{test} , we want m to be as large as possible so that error of validation set is a good representation of the actual test set
 - Typically m is 20% of training set

6.5.1 Effects of m

- We have 2 opposite cases
 - Recall : Model f is trained on the full training set of n data, and f' is trained on the $n - m$ training set
 - Case #1 : Small number m of validation data / large number $n - m$ of training data
 - Advantage : $L_{\text{train}}(f') + L_{\text{val}}(f')$ as f' is well estimated.
 - Limitation : $L_{\text{train}}(f') + L_{\text{val}}(f')$ as f' is too small.
 - Case #2 : Large number m of validation data / small number $n - m$ of training data
 - Advantage : $L_{\text{train}}(f') + L_{\text{val}}(f')$ as the validation set is large enough.
 - Limitation : $L_{\text{train}}(f') + L_{\text{val}}(f')$ as f' is badly estimated.
- How to reconcile the two cases?

6.5.2 k-fold Cross Validation


- k-fold cross-validation technique :
 - Split the original training set into k parts, i.e. **each fold has n/k data points**.
 - Repeat for all folds** : Train on $k-1$ parts and leave one part out as validation set.
- Advantages
 - Each data in the original training set will be used as a validation data**.
 - For each fold, we have a large training set to train a good learner, i.e. $L_{\text{train}}(f') = L_{\text{val}}(f')$.
 - We also have a good estimate of the validation error by averaging the validation error over all folds, i.e. $L_{\text{val}} = \text{mean}(L_{\text{val}}(f'))$.

6.5.3 Early Stopping

- Cross validation extremely computationally expensive
- To speed things up, we will stop optimization after M number of gradient steps, when the validation error starts increasing, even if optimization has not converged yet

7 Miscellaneous

7.1 Matrix Inverse



7.2 Matrix Dot Product

$$A \cdot B = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{bmatrix} \cdot \begin{bmatrix} 1 & 1 & 1 \\ 2 & 2 & 2 \\ 3 & 3 & 3 \end{bmatrix} = \begin{bmatrix} 26 & 16 & 14 \\ 41 & 26 & 22 \\ 56 & 40 & 34 \end{bmatrix}$$

$$= \begin{bmatrix} 1 \cdot 1 + 2 \cdot 2 + 3 \cdot 3 & 1 \cdot 2 + 2 \cdot 2 + 3 \cdot 3 & 1 \cdot 3 + 2 \cdot 3 + 3 \cdot 3 \\ 4 \cdot 1 + 5 \cdot 2 + 6 \cdot 3 & 4 \cdot 2 + 5 \cdot 2 + 6 \cdot 3 & 4 \cdot 3 + 5 \cdot 3 + 6 \cdot 3 \\ 7 \cdot 1 + 8 \cdot 2 + 9 \cdot 3 & 7 \cdot 2 + 8 \cdot 2 + 9 \cdot 3 & 7 \cdot 3 + 8 \cdot 3 + 9 \cdot 3 \end{bmatrix}$$

$$= \begin{bmatrix} 26 & 16 & 14 \\ 41 & 26 & 22 \\ 56 & 40 & 34 \end{bmatrix}$$

7.3 Differentiation Nonsense

- $\frac{d}{dx}(x) = 1$
- $\frac{d}{dx}(ax) = a$
- $\frac{d}{dx}(x^2) = 2x$
- $\frac{d}{dx}(\sin(x)) = \cos(x)$
- $\frac{d}{dx}(\cos(x)) = -\sin(x)$
- $\frac{d}{dx}(\tan(x)) = \sec^2(x)$
- $\frac{d}{dx}(\cot(x)) = -\csc^2(x)$
- $\frac{d}{dx}(\sec(x)) = \sec(x) \tan(x)$
- $\frac{d}{dx}(\csc(x)) = -\csc(x) \cot(x)$
- $\frac{d}{dx}(x^a) = \frac{1}{a} x^{a-1}$
- $\frac{d}{dx}(e^x) = e^x$
- $\frac{d}{dx}(a^x) = \ln(a) a^x$
- $\frac{d}{dx}(\sin^{-1}(x)) = \frac{1}{\sqrt{1-x^2}}$
- $\frac{d}{dx}(\tan^{-1}(x)) = \frac{1}{1+x^2}$
- $\frac{d}{dx}(\sec^{-1}(x)) = \frac{1}{|x| \sqrt{x^2-1}}$

- (Chain Rule) $\frac{d}{dx} f(g(x)) = f'(g(x)) g'(x)$

- (Product Rule) $\frac{d}{dx} (f(x)g(x)) = f(x)g'(x) + f'(x)g(x)$

- (Quotient Rule) $\frac{d}{dx} \left(\frac{f(x)}{g(x)} \right) = \frac{g(x)f'(x) - f(x)g'(x)}{g(x)^2}$

7.4 Useful Equations

- Hyperbolas** :
 - Vertical Opening: $\frac{(y-k)^2}{a^2} - \frac{(x-h)^2}{b^2} = 1$
 - Horizontal Opening: $\frac{(x-h)^2}{a^2} - \frac{(y-k)^2}{b^2} = 1$
- Ellipsoids** (as per your formula):

$$\frac{(x-h)^2}{a^2} + \frac{(y-k)^2}{b^2} = 1$$
 If it's a three-dimensional ellipsoid, you can include a z-term:

$$\frac{(x-h)^2}{a^2} + \frac{(y-k)^2}{b^2} + \frac{(z-l)^2}{c^2} = 1$$
- Circles** :

$$(x-h)^2 + (y-k)^2 = r^2$$
 Here, (h, k) are the coordinates of the center of the shapes, a, b , and c are the semi-axis lengths, and r is the radius of the circle.

7.5 When is ML not useful?

- the task is fully deterministic
- there's a straightforward formula equating x and y
- the task demands different outputs every time, regardless of inputs (randomness)

7.6 kNN for Regression task

2.1 k-NN for Regression

- k-NNs can be used for regression as well as classification. With a new observation x_{test} , calculate the mean of the k nearest neighbours and assign it as the prediction for this new observation.

$$\hat{y}_{\text{test}} = \frac{1}{k} \sum_{i=1}^k x_i \quad x_i \in S$$

- where S is the sorted distances between all $x \in X$.

7.7 Logistic Regression as Binary Classification

- While Linear Regression is for regression, Logistic Regression is for classification – mostly problems with 2 classes (i.e., binary classification). Logistic Regression is similar to Linear Regression in that we compute $\hat{y} = w^T \cdot x + b$. However, for Logistic Regression, \hat{y} is a real number and not a classification. To make the decision of classifying x as class 0 or 1, we use the Sigmoid Function:

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

It takes any number and squashes it into the range $[0, 1]$. If the final Sigmoid value is ≤ 0.5 , we classify it as class 0, and it's ≥ 0.5 , we classify it as class 1:

$$\hat{y} = \begin{cases} 0 & \sigma(z) < 0.5 \\ 1 & \sigma(z) \geq 0.5 \end{cases}$$

7.8 Types of noise

